

18.337 Parallel Prefix

The Parallel Prefix Method

- This is our first example of a parallel algorithm
- Watch closely what is being optimized for
 - Parallel steps
- Beautiful idea with surprising uses
- Not sure if the parallel prefix method is used much in the real world
 - Might maybe be inside MPI scan
 - Might be used in some SIMD and SIMD like cases
- The real key: What is it about the real world that differs from the naïve mental model of parallelism?

Students early mental models

- Look up or figure out how to do things in parallel
- Then we get speedups!
 - NOT!

Parallel Prefix Algorithms

1. A theoretical (may or may not be practical) secret to turning serial into parallel
2. Suppose you bump into a parallel algorithm that surprises you→
“there is no way to parallelize this algorithm” you say
3. Probably a variation on parallel prefix!

Example of a prefix

Sum Prefix

Input $x = (x_1, x_2, \dots, x_n)$

Output $y = (y_1, y_2, \dots, y_n)$

$$y_i = \sum_{j=1:i} x_j$$

Example

$x = (1, 2, 3, 4, 5, 6, 7, 8)$

$y = (1, 3, 6, 10, 15, 21, 28, 36)$

Prefix Functions-- outputs depend upon an *initial* string

What do you think?

- Can we really parallelize this?
- It looks like this sort of code:
 $y=0;$
 for $i=2:n$, $y(i)=y(i-1)+x(i);$ end
- The i th iteration of the loop is not at all decoupled from the $(i-1)$ st iteration.
- Impossible to parallelize right?

A clue?

$$x = (1, 2, 3, 4, 5, 6, 7, 8)$$

$$y = (1, 3, 6, 10, 15, 21, 28, 36)$$

Is there any value in adding, say, $4+5+6+7$?

Note if we separately have $1+2+3$, what can we do?

Suppose we added $1+2$, $3+4$, etc. pairwise, what could we do?

Prefix Functions -- outputs depend upon an *initial* string

Suffix Functions -- outputs depend upon a *final* string

Other Notations

1.+ \ “plus scan” APL (“A Programming Language” source of the very name “scan”, an array based language that was ahead of its time)


2.MPI_scan

3.MATLAB command: $y = \text{cumsum}(x)$

4.MATLAB matmul: $y = \text{tril}(\text{ones}(n)) * x$

Parallel Prefix Recursive View

prefix([1 2 3 4 5 6 7 8]) = [1 3 6 10 15 21 28 36]



Pairwise sums

Recursive prefix

Update “odds”

- Any associative operator

$$\begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$



```
function prefix!(x, *)  
    n=length(x)  
    if n<=1  
        return()  
    end
```

```
    for i=2:2:n # even: pairwise sums  
        x[i] *= x[i-1]  
    end
```

```
    prefix!(view(x,2:2:n),*) # recursive prefix (in place!)
```

```
    for i = 3:2:n # odd: x[i-1] is a cumsum, x[i] is raw data  
        x[i] *= x[i-1]  
    end  
end
```

Operation Count

- Notice
- # adds = $2n$
- # required = n
- Parallelism at the cost of more work!

Any Associative Operation works

Associative:

$$(a+b)+c = a+(b+c)$$

Sum (+)

Product (*)

Max

Min

Input: Reals

All (=and)

Any (= or)

Input: Bits
(Boolean)

MatMul

Inputs: Matrices

Fibonacci via Matrix Multiply Prefix

$$F_{n+1} = F_n + F_{n-1}$$

$$\begin{pmatrix} F_{n+1} \\ F_n \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} F_n \\ F_{n-1} \end{pmatrix}$$

Can compute all F_n by matmul_prefix on

$$\left[\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}, \begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} \right]$$

then select the upper left entry

Arithmetic Modulo 2 (binary arithmetic)

$0+0=0$ $0+1=1$ $1+0=1$ $1+1=0$	$0*0=0$ $0*1=0$ $1*0=0$ $1*1=1$
Add = exclusive or	Mult = and

Carry-Look Ahead Addition (Babbage 1800' s)

Example						
1	0	1	1	1		Carry
	1	0	1	1	1	First Int
	1	0	1	0	1	Second Int
1	0	1	1	0	0	Sum

Goal: Add Two n-bit Integers

Carry-Look Ahead Addition (Babbage 1800' s)

Goal: Add Two n-bit Integers

Example						Notation					
1	0	1	1	1		c_2	c_1	c_0			
1	0	1	1	1	First Int	a_3	a_2	a_1	a_0		
1	0	1	0	1	Second Int	b_3	b_2	b_1	b_0		

Carry-Look Ahead Addition (Babbage 1800's)

Goal: Add Two n-bit Integers

Example						Notation			
1	0	1	1	1		c_2	c_1	c_0	
1	0	1	1	1	First Int	a_3	a_2	a_1	a_0
1	0	1	0	1	Second Int	a_3	b_2	b_1	b_0

(addition mod 2) $c_{-1} = 0$
 for $i = 0 : n-1$

$$s_i = a_i + b_i + c_{i-1}$$

$$c_i = a_i b_i + c_{i-1}(a_i + b_i)$$

end

$$S_n = c_{n-1}$$

Goal: Add Two n-bit Integers

Example						Notation			
1	0	1	1	1		c_2	c_1	c_0	
1	0	1	1	1	First Int	a_3	a_2	a_1	a_0
1	0	1	0	1	Second Int	a_3	b_2	b_1	b_0

$$c_{-1} = 0$$

(addition mod 2)
for $i = 0 : n-1$

$$s_i = a_i + b_i + c_{i-1}$$

$$c_i = a_i b_i + c_{i-1}(a_i + b_i)$$

$$\begin{bmatrix} c_i \\ 1 \end{bmatrix} = \begin{bmatrix} a_i + b_i & a_i b_i \\ 0 & 1 \end{bmatrix} \begin{bmatrix} c_{i-1} \\ 1 \end{bmatrix}$$

end

Goal: Add Two n-bit Integers

Example						Notation					
1	0	1	1	1		c_2	c_1	c_0			
1	0	1	1	1	First Int	a_3	a_2	a_1	a_0		
1	0	1	0	1	Second Int	a_3	b_2	b_1	b_0		

$c_{-1} = 0$ (addition mod 2)

for $i = 0 : n-1$

$$s_i = a_i + b_i + c_{i-1}$$

$$c_i = a_i b_i + c_{i-1}(a_i + b_i)$$

end

$$\begin{bmatrix} c_i \\ 1 \end{bmatrix} = \begin{bmatrix} a_i + b_i & a_i b_i \\ 0 & 1 \end{bmatrix} \begin{bmatrix} c_{i-1} \\ 1 \end{bmatrix}$$

Matmul prefix with binary arithmetic is equivalent to carry-look ahead!

Compute c_i by prefix, then

$s_i = a_i + b_i + c_{i-1}$ in parallel

Tridiagonal Factor

$$\mathbf{T} = \begin{bmatrix} a_1 & b_1 & & & \\ c_1 & a_2 & b_2 & & \\ & c_2 & a_3 & b_3 & \\ & & c_3 & a_4 & b_4 \\ & & & c_4 & a_5 \end{bmatrix}$$

Determinants ($D_0=1$, $D_1=a_1$)
(D_k is the det of the $k \times k$ upper left):

$$D_n = a_n D_{n-1} - b_{n-1} c_{n-1} D_{n-2}$$

Compute D_n by matmul_prefix

$$\begin{bmatrix} D_n \\ D_{n-1} \end{bmatrix} = \begin{bmatrix} a_n & -b_{n-1}c_{n-1} \\ 1 & 0 \end{bmatrix} \begin{bmatrix} D_{n-1} \\ D_{n-2} \end{bmatrix}$$

$$\mathbf{T} = \begin{bmatrix} 1 & & & \\ & d_1 & b_1 & l_1 \\ & & & \\ & & & \end{bmatrix}$$

$$d_n = D_n / D_{n-1}$$

$$l_n = c_n / d_n$$

3 embarrassing
Parallel + prefix

The “Myth” of $\log n$

The $\log_2 n$ parallel steps is not the main reason for the usefulness of parallel prefix.

Say $n = 1000p$ (1000 summands per processor)

Time = (2000 adds) + ($\log_2 P$ message passings)

fast & embarrassingly parallel

(2000 local adds are serial for each processor of course)

10, 000 adds + 3
communication hops
total speed is as if there
is no communication

Myth of
 $\log n$
Example

80, 000

40, 000

20, 000

10, 000

1 2 3 4 5 6 7 8

$\log_2 n$ = number of steps to add n numbers (NO!!)

Any Prefix
Operation May
Be
Segmented!

Segmented Operations

Inputs = Ordered Pairs

(operand, boolean)

Change of
segment indicated
by switching T/F

e.g. (x, T) or (x, F)

$+_2$	(y, T)	(y, F)
(x, T)	(x+y, T)	(y, F)
(x, F)	(y, T)	(x+y, F)

e. g.	1	2	3	4	5	6	7	8
	T	T	F	F	F	T	F	T
Result	1	3	3	7	12	6	7	8

Copy Prefix: $\mathbf{x} + \mathbf{y} = \mathbf{x}$

(is associative)

Segmented

1	2		3	4	5		6		7		8
T	T		F	F	F		T		F		T
1	1		3	3	3		6		7		8

High Performance Fortran

	1	2	3	4	5		T	T	T	T	T	-	
A =	6	7	8	9	10		M =	F	F	T	T	T	
	11	12	13	14	15			T	F	T	F	F	-

			1	20	42	67	45
SUM_PREFIX(A) =		7	27	50	76	105	
		18	39	63	90	120	

SUM_SUFFIX(A)			1	3	6	10	15
SUM_PREFIX(A, DIM = 2) =		6	13	21	30	40	
		11	23	36			

			1	14	17	.
SUM_PREFIX(A, MASK = M) =		1	14	25	.	
		12	14	38		

More HPF

Segmented

A =

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15

T T | F | T T | F F

S =

T	T	F	F	F
F	T	T	F	F
T	T	T	T	T

Sum_Prefix (A, SEGMENTS = S)

1	13	3
6	20	
11	32	

Example of Exclusive

A = 1 2 3 4 5

Sum_Prefix(A) 1 3 6 10 15

Sum_Prefix(A, EXCLUSIVE = TRUE)
0 1 3 6 10

(Exclusive: Don't count myself)

Parallel Prefix

`prefix([1 2 3 4 5 6 7 8])=[1 3 6 10 15 21 28 36]`

1 2 3 4 5 6 7 8

Pairwise sums

3 7 11 15

Recursive prefix

3 10 21 36


Update “evens”

1 3 6 10 15 21 28 36

- Any operator
- AKA: `+\` (APL), `cumsum` (Matlab), `MPI_SCAN`,


1 0 0
1 1 0
1 1 1

Variations on Prefix



Variations on Prefix

exclusive([1 2 3 4 5 6 7 8])=[0 1 3 6 10 15 21 28]




The Family...

Directions	Inclusive Exc=0	Exclusive Exc=1	
Left	Prefix	Exc Prefix	

Variations on Prefix

exclusive([1 2 3 4 5 6 7 8])=[0 1 3 6 10 15 21 28]

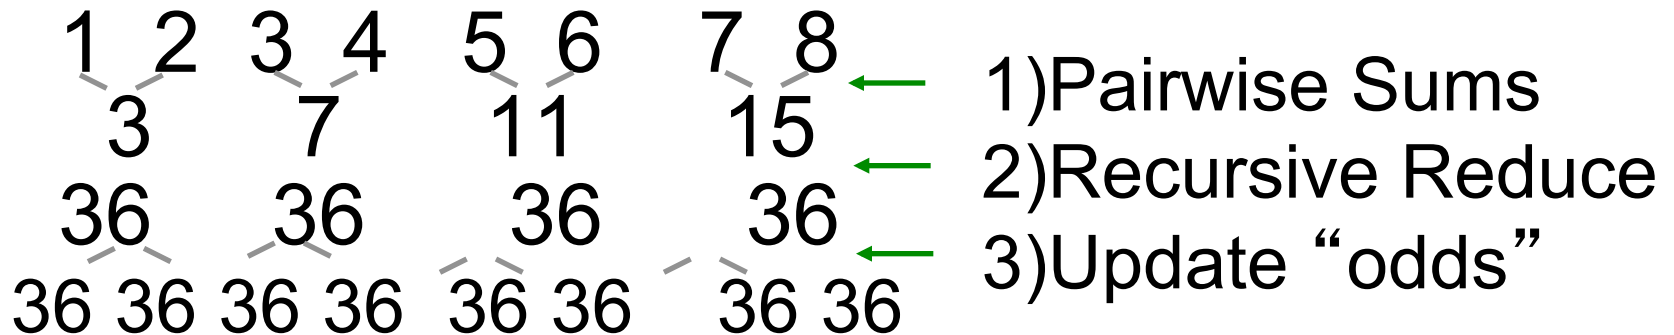


The Family...

Directions	Inclusive Exc=0	Exclusive Exc=1	
Left	Prefix	Exc Prefix	
Right	Suffix	Exc Suffix	

Variations on Prefix

reduce([1 2 3 4 5 6 7 8])=[36 36 36 36 36 36 36 36]




The Family...

Directions	Inclusive Exc=0	Exclusive Exc=1	
Left	Prefix	Exc Prefix	
Right	Suffix	Exc Suffix	
Left/Right	Reduce	Exc Reduce	

Variations on Prefix

exclusive([1 2 3 4 5 6 7 8])=[0 1 3 6 10 15 21 28]



The Family...

Directions	Inclusive Exc=0	Exclusive Exc=1	Neighbor Exc Exc=2
Left	Prefix	Exc Prefix	Left Multipole
Right	Suffix	Exc Suffix	Right " " "
Left/Right	Reduce	Exc Reduce	Multipole

Multipole in 2d or 3d etc

Notice that left/right generalizes more readily to higher dimensions
Ask yourself what Exc=2 looks like in 3d

The Family...

Directions	Inclusive Exc=0	Exclusive Exc=1	Neighbor Exc Exc=2
Left	Prefix	Exc Prefix	Left Multipole
Right	Suffix	Exc Suffix	Right " " "
Left/Right	Reduce	Exc Reduce	Multipole

Not Parallel Prefix but PRAM

- Only concerned with minimizing parallel time (not communication)
- Arbitrary number of processors
- One element per processor

Csanky' s (1977) Matrix Inversion

Lemma 1: (\triangle^{-1}) in $O(\log^2 n)$ (triangular matrix inv)

Proof Idea:
$$\begin{bmatrix} A & 0 \\ C & B \end{bmatrix}^{-1} = \begin{bmatrix} A^{-1} & 0 \\ -B^{-1}CA^{-1} & B^{-1} \end{bmatrix}$$

Lemma 2: *Cayley - Hamilton*

$$p(x) = \det(xI - A) = x^n + c_1 x^{n-1} + \dots + c_n$$

($c_n = \det A$) ±

$$0 = p(A) = A^n + c_1 A^{n-1} + \dots + c_n I$$

$$A^{-1} = (A^{n-1} + c_1 A^{n-2} + \dots + c_{n-1})(-1/c_n)$$

Lemma 3: *Leverier's Lemma*

$$\begin{bmatrix} 1 & & & & \\ s_1 & 2 & & & \\ s_2 & s_1 & \cdot & & \\ \vdots & \vdots & \cdot & \cdot & \\ s_{n-1} & \cdot & \cdot & s_1 & n \end{bmatrix} \begin{bmatrix} c_1 \\ c_2 \\ c_3 \\ \vdots \\ c_n \end{bmatrix} = - \begin{bmatrix} s_1 \\ s_2 \\ s_3 \\ \vdots \\ s_n \end{bmatrix} \quad s_k = \text{tr}(A^k)$$

Csanky

- 1) Parallel Prefix powers of A
- 2) s_k by directly adding diagonals
- 3) c_i from lemmas 1 and 3
- 4) A^{-1} obtained from lemma 2

Matrix multiply can be done in $\log n$ steps
on n^3 processors with the pram model

Can be useful to think this way, but must
also remember how real machines are
built!

- Parallel steps are not the whole story
- Nobody puts one element per processor