

Implementing a parallel Benchmark in Julia

Two versions of PageRank

Outline

- Background
- The PageRank benchmark
- Two Julia implementations
- Results

Background



What are we benchmarking?

- Hardware differences?
- Ability to exploit special hardware features, e.g. fast interconnect?
- Performance of programming language?
- Performance of libraries?
- Efficiency of algorithm?
- Quality of benchmark implementation?

Classes of benchmarks

- Goal oriented
- Algorithm oriented
- Code oriented

The PageRank benchmark



What is PageRank?

- A model for importance of web pages
- 1996: Google founders Larry Page and Sergey Brin for ranking web pages
- 1970s: Very similar work in bibliometrics
- Mathematical structure: Stationary distribution of a Markov chain
- Algorithmic structure:
 - Power iteration on the eigenvalue problem (notice $\lambda = 1$ known)
$$(\alpha P + (1 - \alpha)ve^T)x = x$$
 - Richardson method on linear system
$$(I - \alpha P)x = (1 - \alpha)v$$
- See Gleich (2015)

Why a PageRank Benchmark

- Interest in large data problems (in contrast to physics simulations)
- Mathematically simple
- Exercises a range of (parallel) code pattern
 - Graph generation (local)
 - Sorting (global)
 - Filtering (global)
 - Linear algebra (global)
 - I/O (local)

The PageRank Benchmark Pipeline

- Kernel 0: *Graph generation* (+ write to disk)
- Kernel 1: (Read from disk +) *Sort* (+ write to disk)
- Kernel 2: (Generate adjacency matrix +) *Filter*
- Kernel 3: *PageRank*

Kronecker Graph generation I

- Graphs represented by their adjacency matrices
- Random matrix X is a 2×2 with a *single* unit element (multinomial)

$$P(X_{ij} = 1) = p_{ij} \quad \text{for } i, j = 1, 2$$

- Each edge is

$$E = X^{(1)} \otimes \cdots \otimes X^{(n)} = \bigotimes_{i=1}^n X^{(i)}$$

- Final graph is

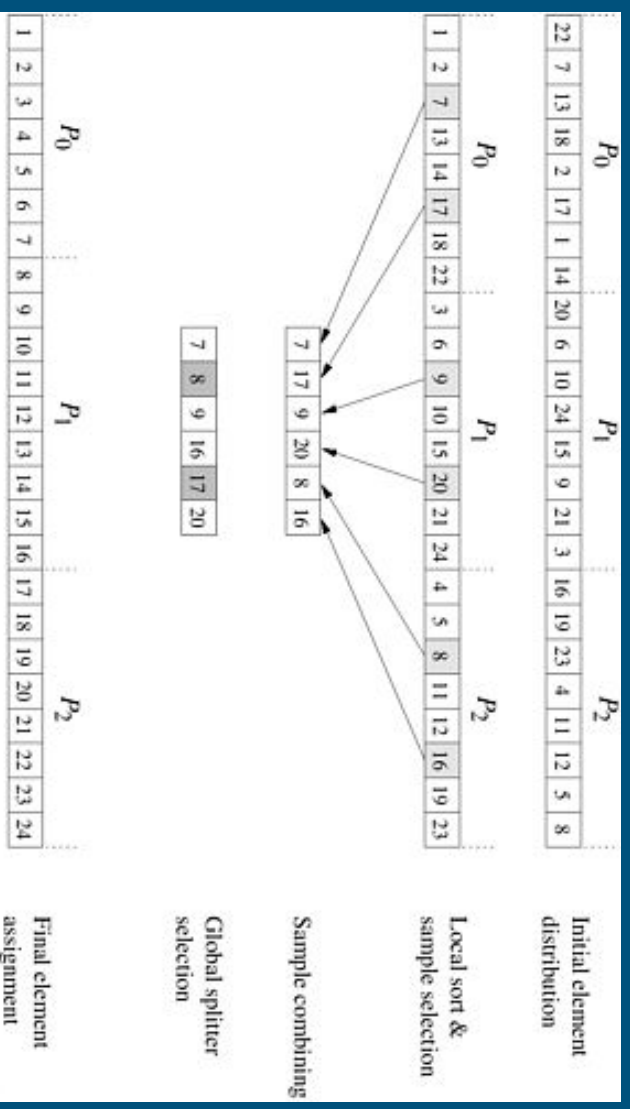
$$G = \sum_{j=1}^m E^{(j)}$$

Kronecker Graph generation II

- Critical to exploit structure
- Data structure should only store list of edges
- E.g. $((a_1, b_1), \dots, (a_n, b_n))$
Or $(a_1, \dots, a_n), (b_1, \dots, b_n)$
- Computations should only track indices of non-zero elements

Sample sort

- Sort edge list by start vertex
- Dominated by all-to-all step



Filtering

- Construct sparse matrix representation of graph (e.g. DCSC)
- Remove edges starting at vertex with highest in-degree

$$\left\{ (i, j) \mid \sum_i G'_{ij} = \max_j \sum_i G'_{ij} \right\}$$

- Remove edges starting at vertices with unit in-degree

$$\left\{ (i, j) \mid \sum_i G'_{ij} = 1 \right\}$$

- Scale rows

$$G_{ij} := \frac{G'_{ij}}{\sum_j G'_{ij}}$$

PageRank

Run 20 iterations of

$$x^{(k+1)} = \alpha P x^{(k)} + (1 - \alpha)v$$

where v is constant

Two Julia implementations



DArrays + remote_call

- Client-Server model
- Distributed/Global Arrays/Partitioned Global Address space (PGAS)
- Mimics semantics of `Array`
- Internally composed by `remote_calls`
 - Similar to Remote procedure call (RPC) / Remote method invocation (RMI)
- Very flexible -> almost all of Julia's datatypes
- Potential overhead

MPI.jl

- SPMD model
- Julia package wrapper C implementations of MPI
- MPI is a message passing model but also
 - Optimized parallel constructs (all-to-all, gather, scatter, broadcast)
 - Optimized transport layers (shared memory, InfinyBand)
- No global abstraction layer for arrays (e.g. global array/PGAS)
- Fast for bittypes

Kronecker graph generation in Julia (common)

```
function Base.rand(x::Kronecker)
    n = 2^x.scl
    m = x.edgesPerVertex*n

    rand2x2(_) = rand(x.onenz)
    randnxn(_) = mapreduce(rand2x2, ⊗, 1:x.scl)
    mapreduce(randnxn, add!, 1:m)
end

rand(Kronecker(scl, edgesPerVertex))
```

Kronecker graph generation in C++

```
for (T i = 0; i < M; i++) {
    T ij_i = 0, ij_j = 0;
    for (int ib = 0; ib < SCALE; ib++) {
        T r1 = random();
        T r2 = random();
        T ii_bit = r1 > ab_scaled;
        T jj_bit = r2 > (c_norm_scaled * ii_bit + a_norm_scaled * !ii_bit);
        ij_i += ii_bit << ib;
        ij_j += jj_bit << ib;
    }
    edges.push_back(std::tuple<T, T>(ij_i, ij_j));
}
```

Kernel 0

- DArray + remote_call

```
map(distribute(files)) do fn
  nEdges = EdgesPerWorker + (myid() == lastWorker ? surplus : 0)
  krongraph500(fn, scl, nEdges)
end
```

- MPI.jl

```
id      = MPI.Comm_rank(state)
nworkers = MPI.Comm_size(state)
nEdges   = EdgesPerWorker + (id + 1 == nworkers ? surplus : 0)
krongraph500(files[id + 1], scl, nworkers == id + 1 ? nEdges)
```